

HDF5 DAOS VOL Connector User's Guide

Jordan Henderson

This document aims to be a helpful guide on how to install the HDF5 DAOS VOL connector and use it to leverage the capabilities of the DAOS object storage system within an HDF5 application.

Contents

1	Introduction	4
2	Installation	5
3	Usage	6
3.1	Using the DAOS VOL connector with an HDF5 application	6
3.2	HDF5 API support	6
3.2.1	H5A interface	6
3.2.2	H5D interface	7
3.2.3	H5F interface	8
3.2.4	H5G interface	9
3.2.5	H5L interface	9
3.2.6	H5O interface	10
3.2.7	H5R interface	11
3.2.8	H5T interface	11
4	Testing	12
4.1	HDF5 and dynamically-loaded VOL connectors	12
4.2	Generic VOL connector test suite	12
5	Appendix	13
5.1	Appendix A - Native VOL connector API calls	13
5.1.1	H5A interface	13
5.1.2	H5D interface	13
5.1.3	H5F interface	14

5.1.4	H5G interface	14
5.1.5	H5L interface	15
5.1.6	H5O interface	15
5.1.7	H5R interface	15
5.1.8	H5T interface	15
6	Revision History	16

1 Introduction

The HDF5 DAOS VOL connector is a connector for HDF5 designed with the goal of allowing HDF5 applications to utilize the DAOS object storage system by translating HDF5 API calls into DAOS calls, as defined by the DAOS API. The DAOS VOL connector is currently built as a dynamically-loaded library that is external to HDF5 and is treated similar to dynamically-loaded HDF5 filter plugins.

For the latest instructions on installation and usage of the HDF5 DAOS VOL connector, please refer to the README file in the connector's source code repository¹; the instructions are listed here for convenience.

¹DAOS VOL Connector README

2 Installation

TBD

3 Usage

3.1 Using the DAOS VOL connector with an HDF5 application

TBD

3.2 HDF5 API support

The following sections serve to illustrate the DAOS VOL connector's support for the HDF5 API, as well as to highlight any differences between the expected behavior of an HDF5 API call versus the actual behavior as implemented by the VOL connector.

3.2.1 H5A interface

Supported API calls

API call	Notes
H5Acreate(1/2)	
H5Acreate_by_name	
H5Aopen(_by_name)	
H5Aopen_name	
H5Awrite	
H5Aread	
H5Aclose	
H5Aiterate(2)	<ul style="list-style-type: none"> Restarting iteration from an index value is currently unsupported Only H5_ITER_NATIVE is supported for the iteration order
H5Aiterate_by_name	<ul style="list-style-type: none"> Restarting iteration from an index value is currently unsupported Only H5_ITER_NATIVE is supported for the iteration order
H5Aget_name	
H5Aget_space	
H5Aget_type	

Currently unsupported API calls

API call	Notes
H5Aopen_by_idx	Currently no support for attribute creation order
H5Aopen_idx	Currently no support for attribute creation order
H5Aexists(_by_name)	
H5Arename(_by_name)	
H5Adelete(_by_name/_by_idx)	
H5Aget_info(_by_name/_by_idx)	
H5Aget_name_by_idx	Currently no support for attribute creation order
H5Aget_create_plist	Default ACPL is currently returned
H5Aget_storage_size	

3.2.2 H5D interface

Supported API calls

API call	Notes
H5Dcreate(1/2)	
H5Dcreate_anon	
H5Dopen(1/2)	
H5Dwrite	
H5Dread	
H5Dclose	
H5Diterate	
H5Dget_space	
H5Dget_type	
H5Dget_create_plist	
H5Dget_access_plist	

Currently unsupported API calls

API call	Notes
H5Dflush	
H5Drefresh	
H5Dextend	
H5Dset_extent	
H5Dget_storage_size	
H5Dget_space_status	Space status is currently always set to H5D_SPACE_STATUS_NOT_ALLOCATED
H5Dget_offset	
H5Dvlen_reclaim	
H5Dvlen_get_buf_size	
H5Dscatter	
H5Dgather	
H5Dfill	

3.2.3 H5F interface

Supported API calls

API call	Notes
H5Fcreate	<ul style="list-style-type: none"> H5F_ACC_EXCL flag is currently not handled correctly
H5Fopen	
H5Freopen	
H5Fis_accessible	
H5Fget_create_plist	
H5Fget_access_plist	
H5Fget_intent	
H5Fget_name	
H5Fclose	

Currently unsupported API calls

API call	Notes
H5Fflush	
H5Fmount	
H5Funmount	
H5Fget_obj_count	
H5Fget_obj_ids	

3.2.4 H5G interface

Supported API calls

API call	Notes
H5Gcreate(1/2)	
H5Gcreate_anon	
H5Gopen(1/2)	
H5Gclose	
H5Gget_create_plist	
H5Gget_info(_by_name/_by_idx)	<p>Of the four fields in the <code>H5G_info_t</code> struct:</p> <ul style="list-style-type: none"> • <code>storage_type</code> is always set to <code>H5G_STORAGE_TYPE_UNKNOWN</code> • <code>nlinks</code> is set appropriately • <code>max_corder</code> is currently always set to 0 • <code>mounted</code> is currently always set to <code>FALSE</code> <p><code>H5Gget_info_by_idx</code> is currently unsupported due to the lack of support for link creation order</p>
H5Gget_num_objs	
H5Glink(2)	Currently only soft link creation is supported

Currently unsupported API calls

API call	Notes
H5Gflush	
H5Grefresh	
H5Gmove(2)	Will be supported once <code>H5Lmove</code> is supported
H5Gunlink	Will be supported once <code>H5Ldelete</code> is supported
H5Gget_linkval	Will be supported once <code>H5Lget_val</code> is supported
H5Gget_objname_by_idx	Will be supported once <code>H5Lget_name_by_idx</code> is supported

3.2.5 H5L interface

Supported API calls

API call	Notes
H5Lcreate_soft	
H5Lexists	
H5Literate(_by_name)	<ul style="list-style-type: none"> Restarting iteration from an index value is currently unsupported Only H5_ITER_NATIVE is supported for the iteration order

Currently unsupported API calls

API call	Notes
H5Lcopy	
H5Lmove	
H5Lcreate_hard	
H5Lcreate_external	
H5Lcreate_ud	
H5Ldelete(_by_idx)	
H5Lget_info(_by_idx)	
H5Lget_name_by_idx	Currently no support for link creation order
H5Lget_val(_by_idx)	
H5Lvisit(_by_name)	
H5Lregister	
H5Lunregister	
H5Lis_registered	
H5Lunpack_elink_val	

3.2.6 H5O interface

Supported API calls

API call	Notes
H5Oopen	
H5Oopen_by_addr	
H5Oclose	

Currently unsupported API calls

API call	Notes
H5Oopen_by_idx	
H5Oexists_by_name	
H5Ovisit(1/2)	
H5Ovisit_by_name(1/2)	
H5Olink	
H5Ocopy	
H5Oflush	
H5Orefresh	
H5Oincr_refcount	
H5Odecr_refcount	

3.2.7 H5R interface

Supported API calls

API call	Notes

Currently unsupported API calls

API call	Notes
H5Rcreate	
H5Rdereference(1/2)	Causes an assertion failure
H5Rget_name	
H5Rget_obj_type(1/2)	
H5Rget_region	

3.2.8 H5T interface

Supported API calls

API call	Notes
H5Tcommit(1/2)	
H5Tcommit_anon	
H5Topen(1/2)	
H5Tclose	
H5Tget_create_plist	

Currently unsupported API calls

API call	Notes
H5Tflush	
H5Trefresh	

4 Testing

4.1 HDF5 and dynamically-loaded VOL connectors

HDF5 has the capability to dynamically load and use a VOL connector for running tests with. While several HDF5 tests have been updated to take advantage of this capability, please be aware that many of these tests are likely to fail or crash due to their native HDF5-specific nature.

In order to choose a particular VOL connector to use for testing, two initial steps must be taken. First, one must help HDF5 locate the VOL connector by pointing to the directory which contains the built library. This can be accomplished by setting the environment variable `HDF5_PLUGIN_PATH` to this directory. Next, HDF5 needs to know the name of which library to use, which is configured by setting the environment variable `HDF5_VOL_CONNECTOR` to the name of the connector.

In order to use the DAOS VOL connector, the aforementioned environment variables should be set as:

```
HDF5_PLUGIN_PATH=/daos/vol/installation/directory/lib
HDF5_VOL_CONNECTOR=daos
```

Having completed this step, HDF5 will be setup to load the DAOS VOL connector and use it for testing.

4.2 Generic VOL connector test suite

In order to test VOL connectors to make sure that they are functioning as expected, a suite of tests which only use the public HDF5 API has been written. If the DAOS VOL connector was built using the HDF5 source code included in the main repository, this suite of tests will be available under the path:

```
src/hdf5/test/VOL/vol_test
```

Currently, this test suite does not have the capability to query what kind of functionality a VOL connector supports and therefore a test will fail if it uses an HDF5 API call which is not implemented, or which is specifically unsupported, in a given VOL connector.

5 Appendix

5.1 Appendix A - Native VOL connector API calls

The following HDF5 API calls are specific to the native HDF5 VOL connector and thus are not able to be implemented by the DAOS VOL connector (or other VOL connectors):

5.1.1 H5A interface

API call	Notes
H5Aiterate1	Deprecated in favor of H5Aiterate2
H5Aget_num_attrs	Deprecated in favor of H5Oget_info

5.1.2 H5D interface

API call	Notes
H5Dformat_convert	
H5Dget_chunk_index_type	
H5Dget_chunk_storage_size	
H5Dread_chunk	
H5Dwrite_chunk	

5.1.3 H5F interface

API call	Notes
H5Fis_hdf5	Uses a default FAPL so can only ever be routed through the native HDF5 VOL connector
H5Fget_vfd_handle	
H5Fget_freespace	
H5Fget_filesize	
H5Fget_file_image	
H5Fget_mdc_config	
H5Fset_mdc_config	
H5Fget_mdc_hit_rate	
H5Fget_mdc_size	
H5Freset_mdc_hit_rate_stats	
H5Fget_info(1/2)	
H5Fget_metadata_read_retry_info	
H5Fget_free_sections	
H5Fclear_elink_file_cache	
H5Fstart_swmr_write	
H5Fstart_mdc_logging	
H5Fstop_mdc_logging	
H5Fget_mdc_logging_status	
H5Fset_libver_bounds	
H5Fformat_convert	
H5Freset_page_buffering_stats	
H5Fget_page_buffering_stats	
H5Fget_mdc_image_info	
H5Fget_eoa	
H5Fincrement_filesize	
H5Fget_dset_no_attrs_hint	
H5Fset_dset_no_attrs_hint	
H5Fset_latest_format	
H5Fset_mpi_atomicity	
H5Fget_mpi_atomicity	

5.1.4 H5G interface

API call	Notes
H5Gset_comment	Deprecated in favor of H5Oset_comment/H5Oset_comment_by_name
H5Gget_comment	Deprecated in favor of H5Oget_comment/H5Oget_comment_by_name
H5Giterate	Deprecated in favor of H5Literate
H5Gget_objinfo	Deprecated in favor of H5Lget_info/H5Oget_info
H5Gget_objtype_by_idx	Deprecated in favor of H5Lget_info/H5Oget_info

5.1.5 H5L interface

API call	Notes

5.1.6 H5O interface

API call	Notes
H5Oget_info(1/2)	
H5Oget_info_by_name(1/2)	
H5Oget_info_by_idx(1/2)	
H5Oset_comment(_by_name)	Deprecated in favor of using attributes on objects
H5Oget_comment(_by_name)	Deprecated in favor of using attributes on objects
H5Oare_mdc_flushes_disabled	
H5Oenable_mdc_flushes	
H5Odisable_mdc_flushes	

5.1.7 H5R interface

API call	Notes

5.1.8 H5T interface

API call	Notes

6 Revision History

February 28, 2019	First Draft